

Anuj Maharjan

Cleveland, OH | 419.494.7550 | anjmhrjn1@gmail.com | maharjananuj.com.np | linkedin | github
Software Engineer | Data Engineer | AWS & Databricks Certified

SUMMARY

Software Engineer and certified Data Engineer with 3+ years of experience building scalable backend systems, Spark-based lakehouse pipelines on Databricks, and production-grade AI workflows. Proficient in Python, SQL, and cloud-native architecture, with hands-on work in dimensional modeling, medallion pipeline design, and LLM orchestration. Recently completed an MSc in Computer Science (AI) at the University of Toledo, combining applied ML research with a strong foundation in distributed systems and software engineering.

SKILLS

Languages & Frameworks: Python, Java, JavaScript, TypeScript, SQL (PostgreSQL, MySQL), NoSQL (MongoDB, DynamoDB), Node.js, React, Next.js, FastAPI, Django Rest Framework

Data Engineering: Apache Spark (PySpark, Spark SQL), Databricks (Unity Catalog, Workflows, Asset Bundles, Auto Loader), Delta Lake, medallion architecture, dimensional modeling, ETL/ingestion pipelines, query optimization

AI/ML: LangChain, PyTorch, Pandas, Numpy, Scikit-learn, TensorFlow, OpenAI, Claude Code, LLM Fine-Tuning, Retrieval-Augmented Generation, VectorDB

Infrastructure & Cloud: AWS (EC2, Lambda, S3, SageMaker, IAM), Docker, Kubernetes, CI/CD, Git

EDUCATION

The University of Toledo, MSc. Computer Science (concentration in AI)

Toledo, OH (2024 – 2026)

Softwarica College of IT and E-Commerce, BSc. (Hons) in Computing

Kathmandu, Nepal (2019 – 2022)

EXPERIENCE

Software Engineer | DigiTechnology

Aug 2022 – July 2024

- Architected core backend services supporting **1,000+ concurrent users**, designing for fault tolerance and data integrity across mission-critical operations.
- Built the reporting and visualization layer of a learning management system, replacing a third-party reporting tool and saving **\$1,500 in annual licensing fees**.
- Developed an in-house **drag-and-drop document builder** using React, Node.js, and NoSQL, cutting document generation time by **50%** and eliminating external SaaS dependencies.
- Optimized RESTful API and CI/CD workflows in collaboration with cross-functional stakeholders, improving deployment reliability and reducing system response times by **40%**.

Software Developer | Kaii Group

Aug 2021 – July 2022

- Engineered a structured **automated data ingestion pipeline** to extract, normalize, and load domain data, improving workflow efficiency by **30%** and eliminating manual intervention across the data lifecycle.
- Built backend services with **Django Rest Framework** and implemented an **automated regression test suite**, reducing production-level bugs by **25%** through systematic coverage of critical code paths.
- Diagnosed and resolved database performance bottlenecks through query analysis and schema optimization, improving PostgreSQL response times by **35%**.
- Delivered secure **CRM software** with role-based access control and token-based session management.

PROJECTS

Medallion Lakehouse Pipeline on Databricks | github.com/anjmhrjn/dabs-taxi

- Built a bronze/silver/gold lakehouse as a **five-task Databricks Workflow DAG**, enforcing validity rules at silver and modeling gold as a **Delta Lake star schema** with surrogate keys and foreign key constraints.
- Packaged it as a **Databricks Asset Bundle** whose dev and prod targets derive every schema, volume, and trigger from one variable, with **Unity Catalog** governance and incremental **Auto Loader** ingestion.

Agent-Orchestrated Adaptive RAG | github.com/anjmhrjn/agentic-rag

- Built a fully-local agentic RAG system in which an orchestrator routes each query across five retrieval strategies, coordinating query-classifier, decomposer, and answer-evaluator agents with a bounded two-retry reflection loop.
- Ran a dual-dataset ablation showing decomposition lifts **MRR 0.56→0.72** on structured retrieval but collapses it to 0.10 on multi-hop, and that reflection costs up to 6× latency for no reliable quality gain.

CERTIFICATIONS

Databricks Certified Data Engineer Associate | Databricks

(August 2026)

AWS Certified Machine Learning Engineer – Associate | Amazon Web Services

(August 2025)

AWS Certified Developer – Associate | Amazon Web Services

(March 2025)